

Spark Setup Documentation

Abraham Sharum, Isaac Tunchez, Andrew Appleyard
Class: CS 30003 - Distributed Systems

December 11, 2025

Demonstration Guide

Installation Steps

1. Install Docker and Docker Compose.
2. Clone the repository.
3. Create a root .env with at least:
 - FINNHUB_API_KEY=
 - INFLUXDB_TOKEN=
 - INFLUXDB_USERNAME=
 - INFLUXDB_PASSWORD=
 - INFLUXDB_ORG=
 - INFLUXDB_BUCKET=
4. Ensure ports 4000, 7077, 8001, 8080, 8081, 8082, 8086, 8800 are free.

Run the Containers (Exact Commands)

- Build and start: `docker compose up --build`
- Stop: `docker compose down`
- Rebuild a single service (example backend): `docker compose build backend`
- Spark services: master + two workers come up via compose; workers are limited to `SPARK_WORKER_CORES=2` and `SPARK_WORKER_MEMORY=2g`.

Run the CRUD Application

- Backend FastAPI: <http://localhost:4000> (Swagger UI at /docs).
- Frontend dashboard: <http://localhost:8800>.
- Key endpoints:
 - `GET /portfolio` — list all trades.
 - `POST /trades` — create (e.g., `{"symbol": "AAPL", "shares": 5, "buy_price": 180}`).
 - `PUT /trades/{symbol}` — update buy price (e.g., `{"new_price": 185}`).
 - `DELETE /trades/{symbol}` — delete/sell a trade; optional body `{"sale_price": 190}` to compute net gain.
 - `GET /cache` — view in-memory price cache (refreshed every 5s for tracked symbols).
- Finnhub quote API: <http://localhost:8001/quote/{symbol}> (uses `FINNHUB_API_KEY`).
- Grafana (optional dashboards): <http://localhost:3000>.

Reproduce the Custom Example

1. Start services: `docker compose up --build`.

2. Create a trade:

```
curl -X POST http://localhost:4000/trades \
-H "Content-Type: application/json" \
-d '{"symbol": "AAPL", "shares": 5, "buy_price": 180}'
```

3. Update it:

```
curl -X PUT http://localhost:4000/trades/AAPL \
-H "Content-Type: application/json" \
-d '{"new_price": 185}'
```

4. Delete it:

```
curl -X DELETE http://localhost:4000/trades/AAPL
```

5. Verify via `GET /portfolio` or refresh the frontend dashboard.